

Short Lecture: Face Detection: The Viola-Jones Algorithm

Spring 2021

Slides by: Dr Carole Twining

Presented by: Terence Morley

Handouts & Lecture Notes


- Report in Scientific American (June 2014):
“In each study, however, those who wrote out their notes by hand had a stronger conceptual understanding and were more successful in applying and integrating the material than those who used [sic] took notes with their laptops.”

The Pen Is Mightier Than the Keyboard


P. A. Mueller, D. M. Oppenheimer, *Psychological Science*, Vol 25, Issue 6, pp. 1159 – 1168, April-23-2014.

- Handouts are to aid note taking, not a total replacement for note taking
- Podcasts, slides, pdfs etc on BlackBoard

Commercial Applications



Sony Cyber-Shot



- Face **Detection**:

- Red-eye removal, focus and exposure
- Tracking
- Adult/**Child** detection (Sony)

- Face **Recognition**:

- Save person-specific settings
- Identity-Based Image Retrieval



- Cat and Dog **Detection**:

- Focus and exposure
- Take picture when face-on
- Pet recognition as well! (Pentax)



Images from Fujifilm

Clues to Detecting (Human) Faces?






- Basic features and their relationship
- Simple pattern in blurred, grayscale, contrast-adjusted patches
- Dark & light patches, simple horizontal & vertical patterns

Task:

- Scan images, and find regions that contain examples of the object of interest




- Learning Approach: Given positive and negative examples, learn best features and best classifier



Fast Face Detection

Requirements:

- Speed of computation
(15 frames per second)
- ~ 90% - 95% detection rate
- ~ 10^{-5} false positive rate




Approach:

- Limit yourself to frontal, upright faces
- Efficient-to-compute features
- Efficient image representation
- Adaboost for efficient choice of features
- Cascade of classifiers



Rectangle Features: Definition


$$\mathcal{F} = \sum_{\vec{r} \in A} \mathcal{I}(\vec{r}) - \sum_{\vec{r} \in B} \mathcal{I}(\vec{r})$$



- Vertical/Horizontal/Diagonal feature detectors
- Implementation as convolution too inefficient

Image Representation: Integral Image

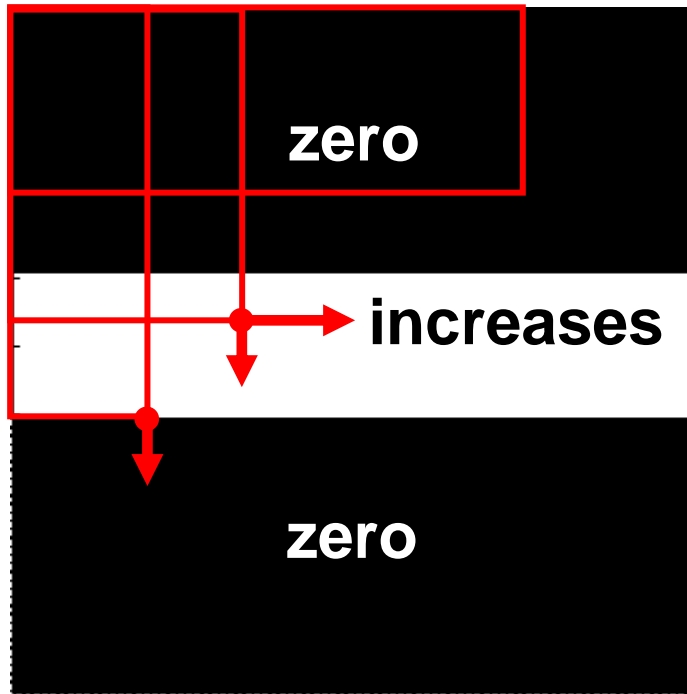
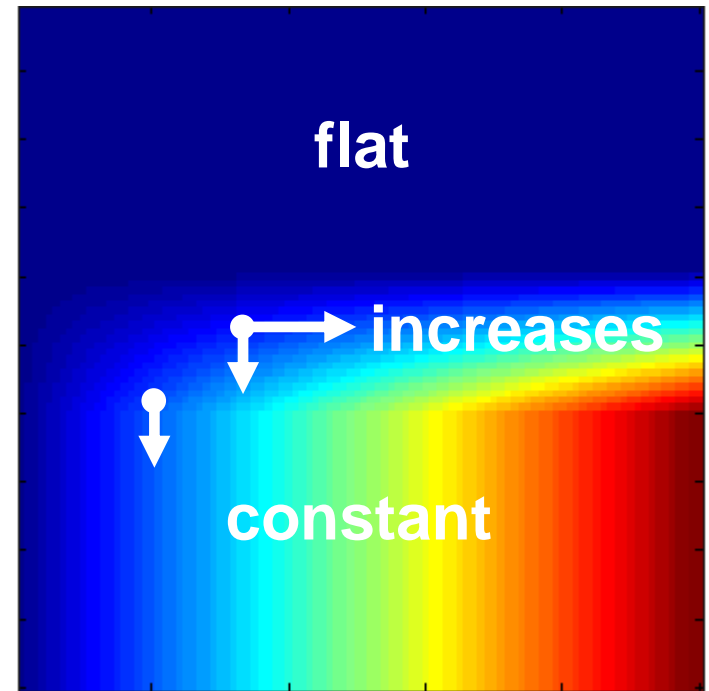
Image



$$\mathcal{I}(x, y) \doteq \sum_{x' \leq x} \sum_{y' \leq y} \mathcal{I}(x', y')$$

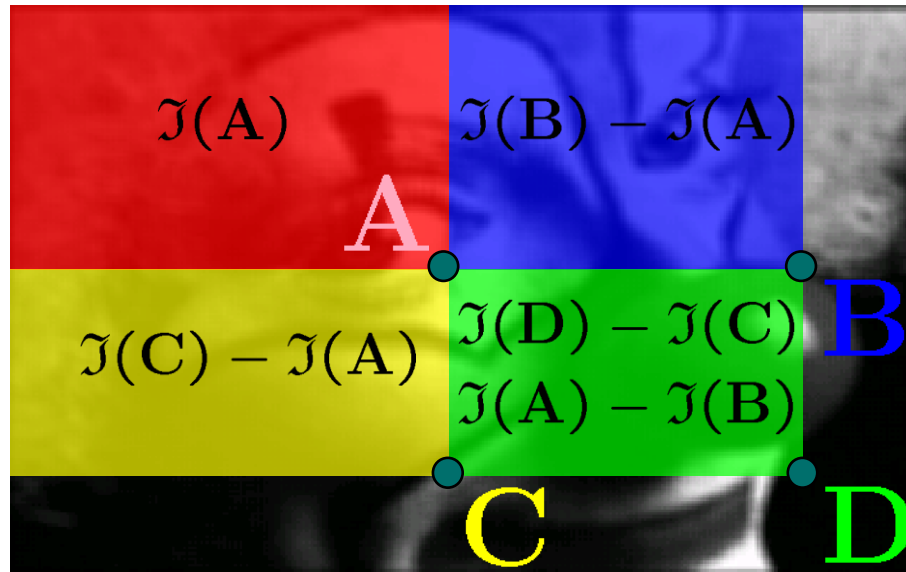
- Integral image can be computed in one pass

Image Representation: Integral Image

 $\mathcal{I}(x, y)$  $\mathcal{J}(x, y)$ 

- Integral image contains all information in original image


Rectangle Features: Integral Image






$$\sum_{\vec{r} \in ABCD} \mathcal{I}(\vec{r}) =$$

$$\mathcal{I}(D) - \mathcal{I}(B) + \mathcal{I}(A) - \mathcal{I}(C)$$


- Single rectangle: 4 array references
- Two-rectangle: 6 array references (<8, shared points)
- Three-rectangle: 8
- Four-rectangle: 9
- Rectangle features efficient using integral image



Rectangle Features on Faces



- 24 x 24 frontal faces, roughly aligned
- 2,3 & 4 rectangle features -> 160,000 in total
- Intensity normalisation using variance:
 - sum of values & sum of values squared
 - integral image of image & integral image of squared image
- Which rectangle features?





AdaBoost Face Classifier

- Training Data:
5000 24 x 24 faces & 10,000 24 x 24 non-faces
- Result: 200 features, 95% detection rate, 1 in 14,084 false positive rate
- AdaBoost first two features:




Eyes darker compared to cheekbones



Eyes darker compared to bridge of nose

Searching the Whole Image:

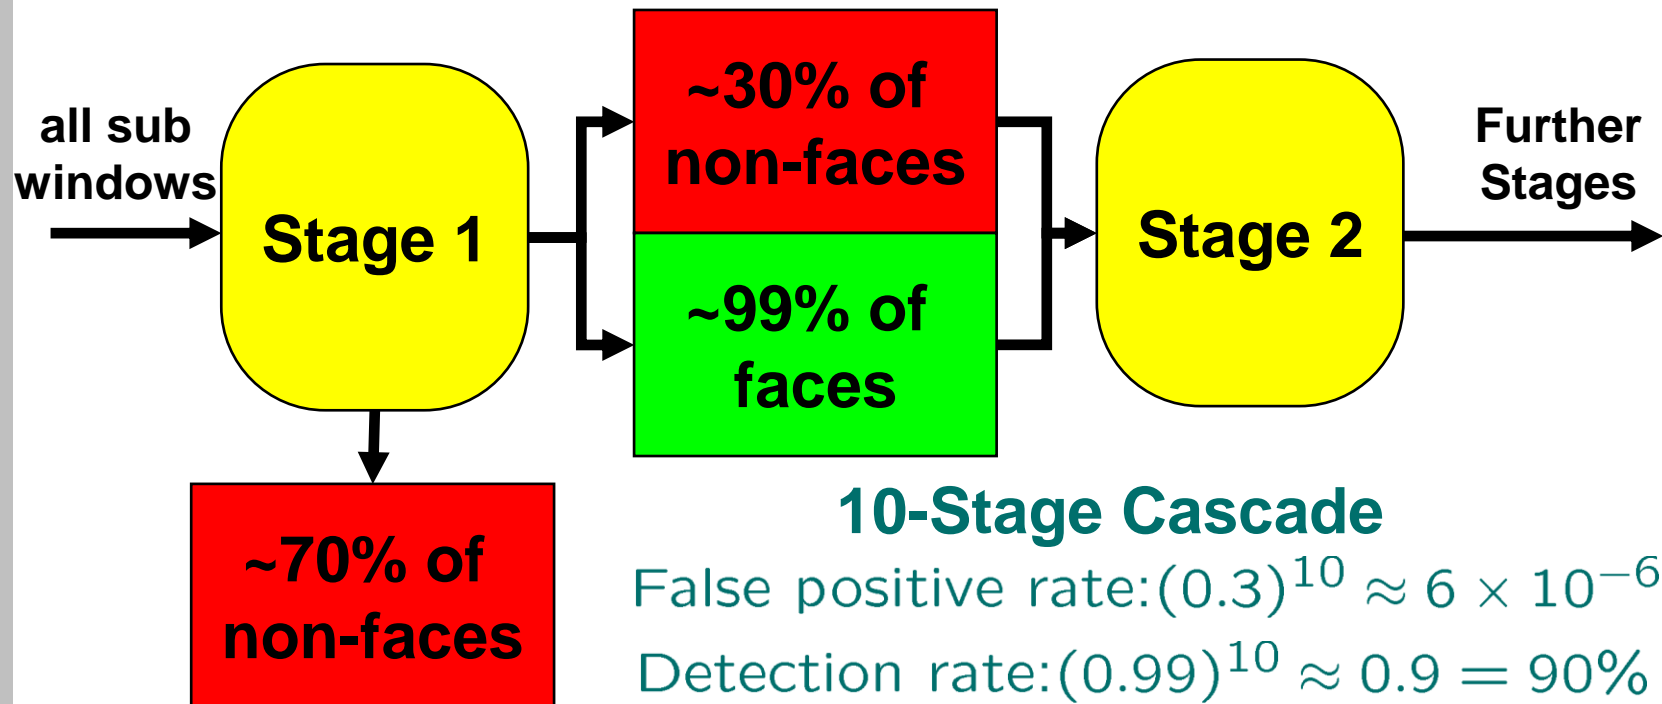
Original Image



- Have to consider a range of possible scales
- Have to consider all possible patch positions
- Most patches NOT faces!

Cascade Classifier

- Single-stage AdaBoost classifier too slow, but shows that rectangle features can detect faces
- In any image, most sub-windows are not faces



**don't waste time on
further processing**


Viola & Jones Cascade Classifier

Paul Viola & Michael Jones, *Robust Real-Time Face Detection*,
International Journal of Computer Vision, 57(2), 137—154 (2004)

- 38 stages of varying complexity, 6060 features in all
- Trained in a day using parallel processing
- Searches 384x288 pixel image in 0.067 seconds
- Applies features at multiple scales and locations
- **Fifteen** times faster than previous approaches
- Can be generalised to other objects
(people/cars/cats/dogs)

Original Face Detection Results:

- Works at various scales and small rotation



Original Face Detection Results:



- Finds multiple faces within a single image

Original Face Detection Results:

- Fails for faces in profile (as expected)
- Occasional mistakes



Online Demo: Upload your own image

Image with Faces Detected




- Tudor Dobrila <http://www.airtudor.com/> **Now defunct!**

Face Detection: Summary

- Fast detectors became available
 - Vary the features (Haar -> SURF etc), Vary classifiers
- Sufficient for commercial applications
- Face detection not the same as facial recognition
 - Identifying a person & their expression
- Apply similar feature detection to detected faces
- Try a different approach: shape models



Training Features



Recording pause



Appendix A: Dogs & Cats

Detecting Cats:

- Cat faces: greater appearance variation than human

Cat Head Detection – Zhang, Sun, and Tang, ECCV 2008

- Use image intensity and image gradient channels




Intensity




Vertical





Horizontal





Two Diagonals



- Within and between channel rectangle/Haar features



Cat Head Detection: Results



Recording pause

Appendix B:

Intensity Normalisation & AdaBoost

Intensity Normalisation:

Mean within window: $\bar{\mathcal{I}} = \frac{1}{N} \sum_{\tilde{\mathbf{r}} \in W} \mathcal{I}(\tilde{\mathbf{r}})$

Variance within window: $\sigma^2 = \frac{1}{N} \sum_{\tilde{\mathbf{r}} \in W} (\mathcal{I}(\tilde{\mathbf{r}}) - \bar{\mathcal{I}})^2$

$$= -(\bar{\mathcal{I}})^2 + \frac{1}{N} \sum_{\tilde{\mathbf{r}} \in W} (\mathcal{I}(\tilde{\mathbf{r}}))^2$$

Intensity normalisation using mean & variance:

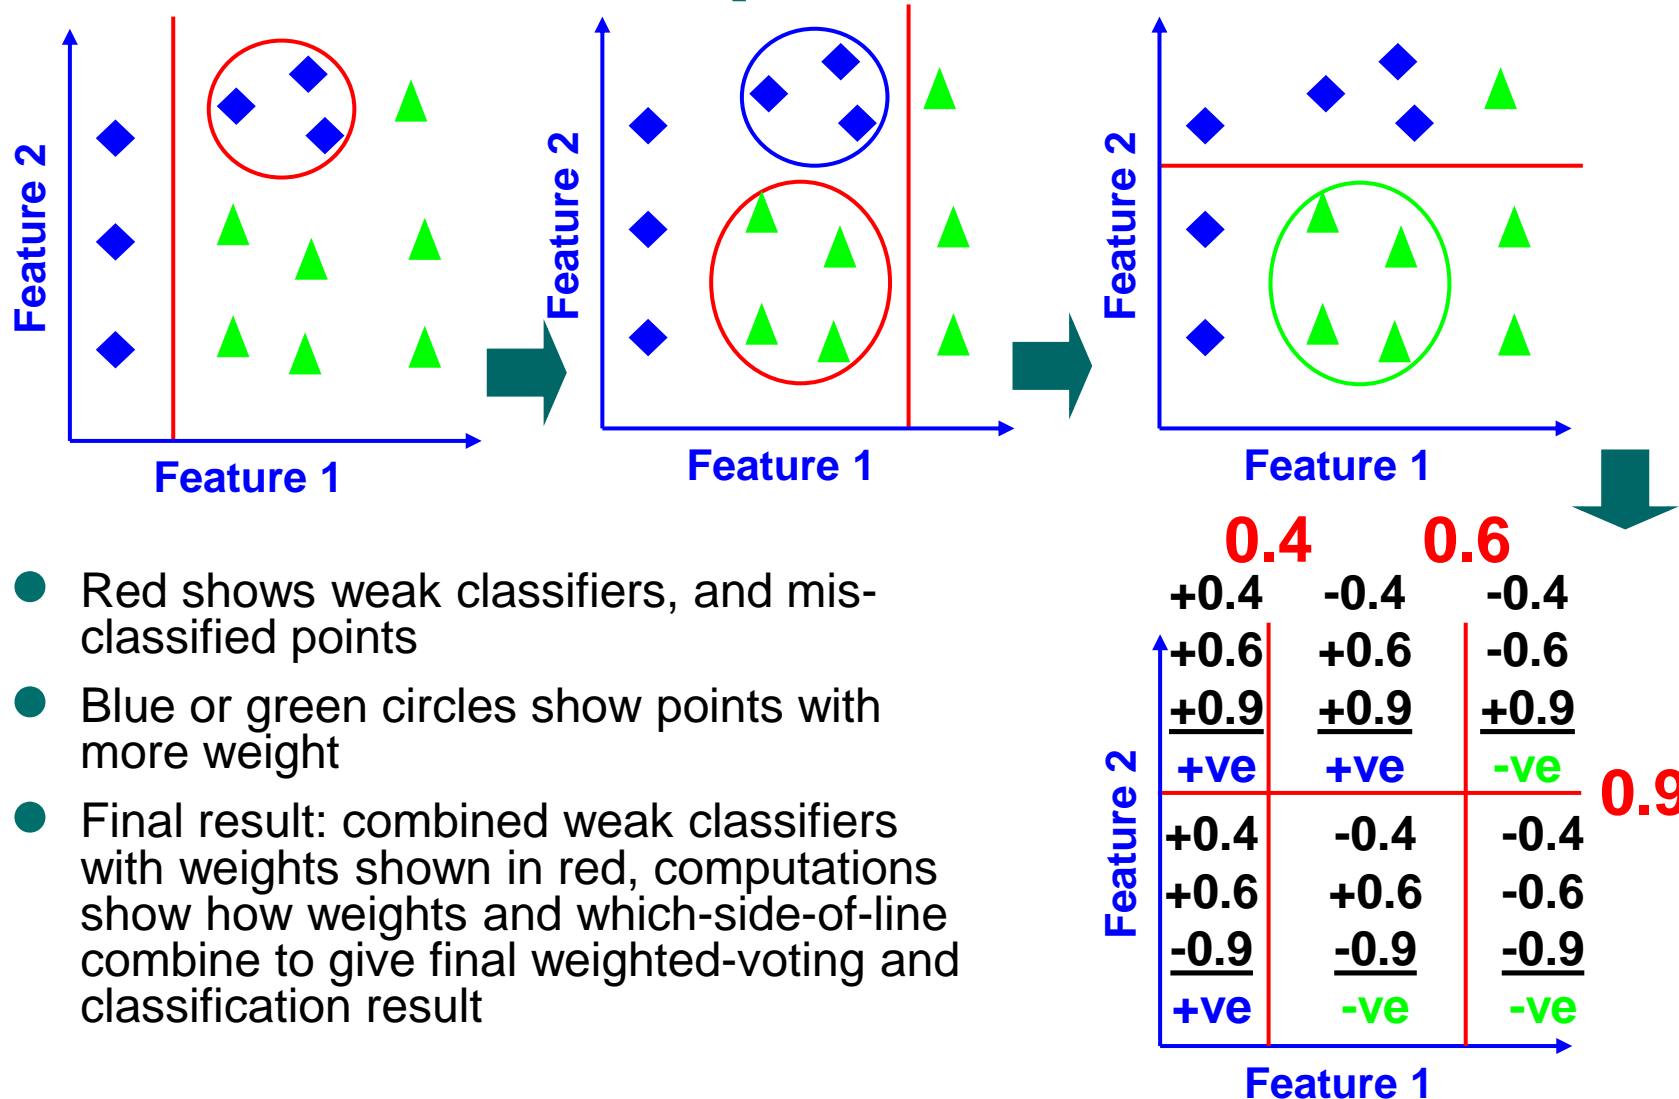
Shift to fixed mean & scale to unit variance

- sum of values
& sum of values squared
- integral image of image
& integral image of squared image

AdaBoost (Adaptive Boosting) Classifier:

- Combination of weak classifiers
- Each weak classifier is based on value of one feature
value < threshold, or value > threshold
- Binary output from each weak classifier: ± 1
- To build: add new weak classifier by concentrating on mis-classified examples at last round
- Output is computed by weighted voting from ALL of the weak classifiers
- Means you have to evaluate ALL of the chosen weak classifiers on each patch to get a decision

Adaboost Example:



- Red shows weak classifiers, and misclassified points
- Blue or green circles show points with more weight
- Final result: combined weak classifiers with weights shown in red, computations show how weights and which-side-of-line combine to give final weighted-voting and classification result